




The Scientific Python Ecosystem

Jim Pivarski

Princeton University – DIANA-HEP


July 25, 2018

Why you can't ignore Python





<http://pypl.github.io/PYPL.html>

Why you can't ignore Python





<http://pypl.github.io/PYPL.html>

Why you can't ignore Python



Why you can't ignore Python



Why you can't ignore Python




All of the deep learning libraries I could find either have a Python interface or are primarily/exclusively Python.



Why you can't ignore Python




Mentions of Software in Astronomy Publications:





Compiled from NASA ADS [\(code\)](#).

Thanks to Juan Nunez-Iglesias,
Thomas P. Robitaille, and Chris Beaumont.

Python's Scientific Stack




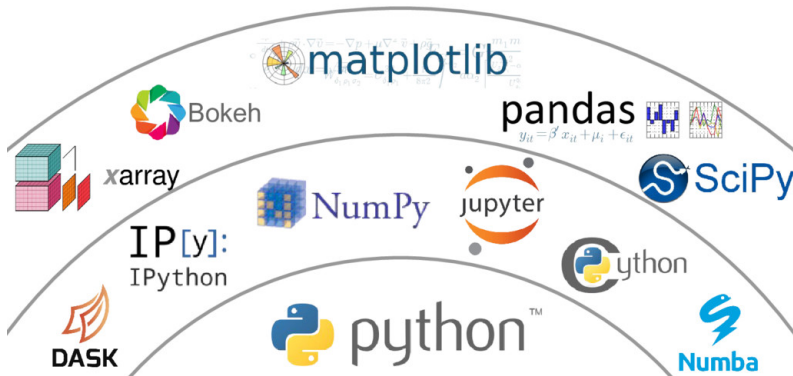
Python's Scientific Stack



More significant: what has grown around Python




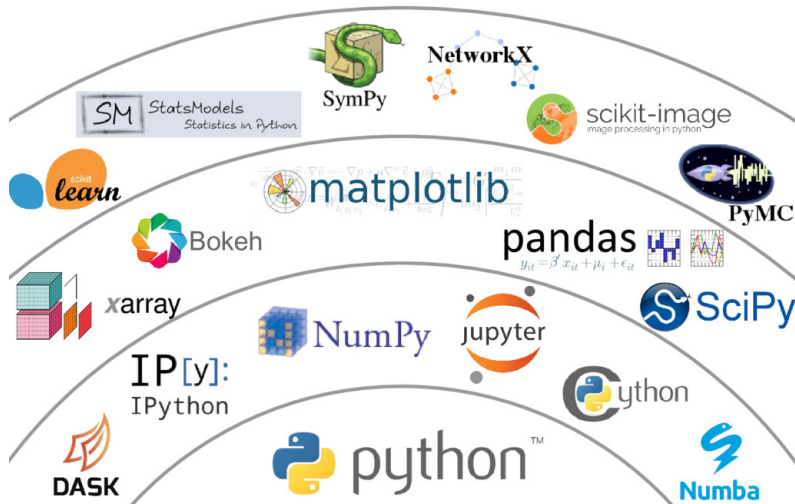
Python's Scientific Stack




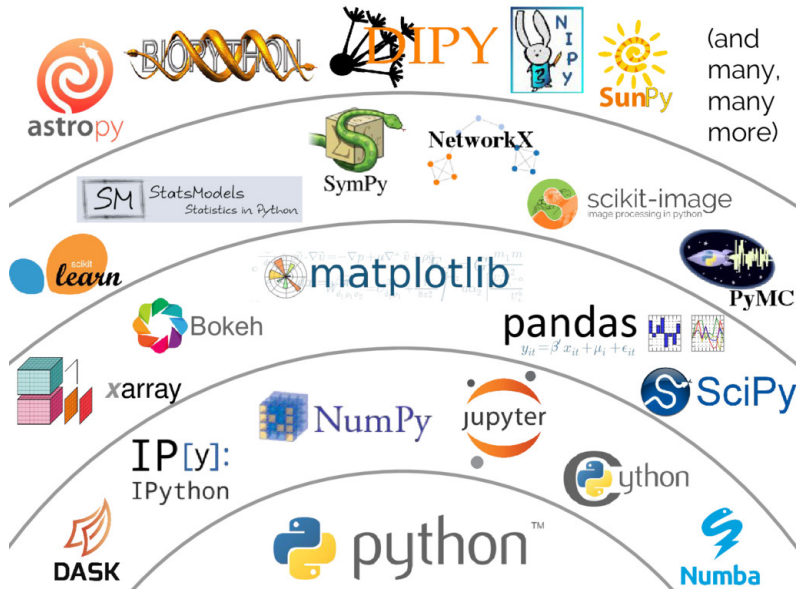
More significant: what has grown around Python



Python's Scientific Stack



More significant: what has grown around Python



Why don't you use C instead of Python? It's so much faster!

Why don't you commute by airplane instead of by car? It's so much faster!



The Unexpected Effectiveness of Python in Science

Jake VanderPlas @jakevdp
PyCon 2017

“Python is a slow language”



Quibble: A language is neither fast nor slow. (But CPython is pretty slow.)

“Python is a slow language”



Quibble: A language is neither fast nor slow. (But CPython is pretty slow.)

Why Python is slow

“Python is a slow language”



Quibble: A language is neither fast nor slow. (But CPython is pretty slow.)

Why Python is slow

- ▶ Virtual machine between Python bytecode and the physical machine.

“Python is a slow language”



Quibble: A language is neither fast nor slow. (But CPython is pretty slow.)

Why Python is slow

- ▶ Virtual machine between Python bytecode and the physical machine.
- ▶ Type-checking during execution, every time in a tight for loop.

“Python is a slow language”



Quibble: A language is neither fast nor slow. (But CPython is pretty slow.)

Why Python is slow

- ▶ Virtual machine between Python bytecode and the physical machine.
- ▶ Type-checking during execution, every time in a tight for loop.
- ▶ Data are bloated and distributed in memory (cache line efficiency, pointer chasing).

“Python is a slow language”



Quibble: A language is neither fast nor slow. (But CPython is pretty slow.)

Why Python is slow

- ▶ Virtual machine between Python bytecode and the physical machine.
- ▶ Type-checking during execution, every time in a tight for loop.
- ▶ Data are bloated and distributed in memory (cache line efficiency, pointer chasing).
- ▶ Global Interpreter Lock (GIL): parallel processing is not parallel.

“Python is a slow language”



Quibble: A language is neither fast nor slow. (But CPython is pretty slow.)

Why Python is slow

- ▶ Virtual machine between Python bytecode and the physical machine.
- ▶ Type-checking during execution, every time in a tight for loop.
- ▶ Data are bloated and distributed in memory (cache line efficiency, pointer chasing).
- ▶ Global Interpreter Lock (GIL): parallel processing is not parallel.

Why Python is nevertheless good for science

“Python is a slow language”



Quibble: A language is neither fast nor slow. (But CPython is pretty slow.)

Why Python is slow

- ▶ Virtual machine between Python bytecode and the physical machine.
- ▶ Type-checking during execution, every time in a tight for loop.
- ▶ Data are bloated and distributed in memory (cache line efficiency, pointer chasing).
- ▶ Global Interpreter Lock (GIL): parallel processing is not parallel.

Why Python is nevertheless good for science

- ▶ Python usually just drives compiled code (good C, C++, Fortran bindings).

“Python is a slow language”



Quibble: A language is neither fast nor slow. (But CPython is pretty slow.)

Why Python is slow

- ▶ Virtual machine between Python bytecode and the physical machine.
- ▶ Type-checking during execution, every time in a tight for loop.
- ▶ Data are bloated and distributed in memory (cache line efficiency, pointer chasing).
- ▶ Global Interpreter Lock (GIL): parallel processing is not parallel.

Why Python is nevertheless good for science

- ▶ Python usually just drives compiled code (good C, C++, Fortran bindings).
- ▶ Many standard precompiled libraries, which also release the GIL.

“Python is a slow language”



Quibble: A language is neither fast nor slow. (But CPython is pretty slow.)

Why Python is slow

- ▶ Virtual machine between Python bytecode and the physical machine.
- ▶ Type-checking during execution, every time in a tight for loop.
- ▶ Data are bloated and distributed in memory (cache line efficiency, pointer chasing).
- ▶ Global Interpreter Lock (GIL): parallel processing is not parallel.

Why Python is nevertheless good for science

- ▶ Python usually just drives compiled code (good C, C++, Fortran bindings).
- ▶ Many standard precompiled libraries, which also release the GIL.
- ▶ **You start thinking in terms of “slow control” and “fast math.”**

“Python is a slow language”



Quibble: A language is neither fast nor slow. (But CPython is pretty slow.)

Why Python is slow

- ▶ Virtual machine between Python bytecode and the physical machine.
- ▶ Type-checking during execution, every time in a tight for loop.
- ▶ Data are bloated and distributed in memory (cache line efficiency, pointer chasing).
- ▶ Global Interpreter Lock (GIL): parallel processing is not parallel.

Why Python is nevertheless good for science

- ▶ Python usually just drives compiled code (good C, C++, Fortran bindings).
- ▶ Many standard precompiled libraries, which also release the GIL.
- ▶ **You start thinking in terms of “slow control” and “fast math.”**
- ▶ **Computational performance is not the same as user productivity!**

Numpy was key to Python developing a scientific ecosystem



- 1994 **Python** 1.0 released.
- 1995 First array package: **Numeric** (a.k.a. Numerical, Numerical Python, NumPy).
- 2001 Diverse scientific codebases merged into **SciPy**.
- 2003 **Matplotlib**
- 2003 Numeric was limited; **numarray** appeared as a competitor with more features (memory-mapped files, alignment, record arrays).
- 2005 Two packages were incompatible; could not integrate numarray-based code into SciPy. Travis Oliphant merged the codebases as **Numpy**.
- 2008 **Pandas**
- 2010 **Scikit-Learn**
- 2011 **AstroPy**
- 2012 **Anaconda**
- 2014 **Jupyter**
- 2015 **Keras**

Numpy was key to Python developing a scientific ecosystem



- 1994 **Python** 1.0 released.
- 1995 First array package: **Numeric** (a.k.a. Numerical, Numerical Python, NumPy).
- 2001 Diverse scientific codebases merged into **SciPy**.
- 2003 **Matplotlib**
- 2003 Numeric was limited; **numarray** appeared as a competitor with more features (memory-mapped files, alignment, record arrays).
- 2005 Two packages were incompatible; could not integrate numarray-based code into SciPy. Travis Oliphant merged the codebases as **Numpy**.
- 2008 **Pandas**
- 2010 **Scikit-Learn**
- 2011 **AstroPy**
- 2012 **Anaconda**
- 2014 **Jupyter**
- 2015 **Keras**

The scientific Python ecosystem could have failed before it started if the Numeric-numarray split hadn't been resolved!



Although you can write Python `for` loops over Numpy arrays, you don't reap the benefit unless you express your calculation in Numpy ufuncs (universal functions).

```
pz = numpy.empty(len(pt))  
for i in range(len(pt)):  
    pz[i] = pt[i]*numpy.sinh(eta[i])
```

VS

```
pz = pt * numpy.sinh(eta)
```

$\mathcal{O}(N)$ Python bytecode instructions, type-checks, interpreter locks.

$\mathcal{O}(1)$ Python bytecode instructions, type-checks, interpreter locks.

$\mathcal{O}(N)$ statically typed, probably vectorized native bytecode operations on contiguous memory.



Although you can write Python `for` loops over Numpy arrays, you don't reap the benefit unless you express your calculation in Numpy ufuncs (universal functions).

```
pz = numpy.empty(len(pt))  
for i in range(len(pt)):  
    pz[i] = pt[i]*numpy.sinh(eta[i])
```

VS `pz = pt * numpy.sinh(eta)`

$\mathcal{O}(N)$ Python bytecode instructions, type-checks, interpreter locks.

$\mathcal{O}(1)$ Python bytecode instructions, type-checks, interpreter locks.

$\mathcal{O}(N)$ statically typed, probably vectorized native bytecode operations on contiguous memory.

In other words, a Single (Python) Instruction on Multiple Data.



Although you can write Python `for` loops over Numpy arrays, you don't reap the benefit unless you express your calculation in Numpy ufuncs (universal functions).

```
pz = numpy.empty(len(pt))  
for i in range(len(pt)):  
    pz[i] = pt[i]*numpy.sinh(eta[i])
```

VS `pz = pt * numpy.sinh(eta)`

$\mathcal{O}(N)$ Python bytecode instructions, type-checks, interpreter locks.


$\mathcal{O}(1)$ Python bytecode instructions, type-checks, interpreter locks.

$\mathcal{O}(N)$ statically typed, probably vectorized native bytecode operations on contiguous memory.

In other words, a Single (Python) Instruction on Multiple Data.

The same code reorganization that speeds up Python with Numpy would speed up C++ with CUDA or SIMD.

The NumPythonic mindset: Numpy API in CUDA



HIGH PERFORMANCE WITH CUDA

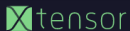
CuPy is an open-source matrix library accelerated with NVIDIA CUDA. It also uses CUDA-related libraries including cuBLAS, cuDNN, cuRand, cuSolver, cuSPARSE, cuFFT and NCCL to make full use of the GPU architecture.



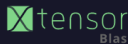
PROJECTS

QuantStack developers contribute to a number of open-source projects, including jupyter, xtensor, bqplot, conda-forge and many others.

High-Performance-Computing



C++ tensor algebra library



BLAS extension to xtensor



C++ wrappers for SIMD intrinsics
and optimized math
implementations



Switch to the notebook now or install packages.

```
# install Miniconda...
```

```
conda install pip
```

```
conda install numpy
```

```
conda install pandas
```

```
# optional...
```

```
conda install numba
```

```
conda install cython
```

```
conda install scipy
```

```
# if you have a GPU and CUDA...
```

```
conda install -c lukepflister pycuda
```